

Training Or Inference Configuration Table

Record reproducibility settings without turning them into performance claims.

Table 1: Training and inference configuration. Values are protocol facts, not performance results.

Phase	Setting	Value or policy
<i>Optimization</i>		
Training	Optimizer	AdamW with cosine decay and no schedule change after model selection.
Training	Batch budget	128 sequences per update with gradient accumulation for memory parity.
<i>Selection</i>		
Validation	Early stopping	Select the checkpoint with the best validation score under the primary metric.
Validation	Random seeds	Report mean and interval over three seeds when a metric is seed-sensitive.
<i>Inference</i>		
Decoding	Budget	Fixed token and tool-call budget across all compared systems.

Width policy: Phase-grouped key-value table; no best-value styling.